

## Data Generation

We'll generate a synthetic dataset that simulates the complex process

```
pip install seaborn
```

Collecting seaborn

Using cached seaborn-0.13.2-py3-none-any.whl.metadata (5.4 kB)

Requirement already satisfied: numpy!=1.24.0,>=1.20 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from seaborn) (1.26.4)

Requirement already satisfied: pandas>=1.2 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from seaborn) (2.2.3)

Requirement already satisfied: matplotlib!=3.6.1,>=3.4 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from seaborn) (3.8.4)

Requirement already satisfied: contourpy>=1.0.1 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.2.1)

Requirement already satisfied: cyclor>=0.10 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (4.51.0)

Requirement already satisfied: kiwisolver>=1.3.1 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (1.4.5)

Requirement already satisfied: packaging>=20.0 in c:\users\yoga\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (23.2)

Requirement already satisfied: pillow>=8 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (10.3.0)

Requirement already satisfied: pyparsing>=2.3.1 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (3.1.2)

Requirement already satisfied: python-dateutil>=2.7 in c:\users\yoga\appdata\roaming\python\python39\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (2.8.2)

Requirement already satisfied: importlib-resources>=3.2.0 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from matplotlib!=3.6.1,>=3.4->seaborn) (6.4.0)

Requirement already satisfied: pytz>=2020.1 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from pandas>=1.2->seaborn) (2024.2)

Requirement already satisfied: tzdata>=2022.7 in c:\users\yoga\appdata\local\programs\python\python39\lib\site-packages (from

```
pandas>=1.2->seaborn) (2024.2)
Requirement already satisfied: zipp>=3.1.0 in c:\users\yoga\appdata\
roaming\python\python39\site-packages (from importlib-
resources>=3.2.0->matplotlib!=3.6.1,>=3.4->seaborn) (3.17.0)
Requirement already satisfied: six>=1.5 in c:\users\yoga\appdata\
roaming\python\python39\site-packages (from python-dateutil>=2.7-
>matplotlib!=3.6.1,>=3.4->seaborn) (1.16.0)
Downloading seaborn-0.13.2-py3-none-any.whl (294 kB)
Installing collected packages: seaborn
Successfully installed seaborn-0.13.2
Note: you may need to restart the kernel to use updated packages.
```

```
[notice] A new release of pip is available: 24.2 -> 24.3.1
[notice] To update, run: python.exe -m pip install --upgrade pip
```

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Set Random Seed To ensure reproducibility.

```
np.random.seed(42)
```

Generate Input Variables Let's assume we have 10 input variables.

```
n_samples = 1000
n_features = 10

# Generate random input features between 0 and 10
X = np.random.uniform(0, 10, size=(n_samples, n_features))
```

Assign Coefficients Randomly assign coefficients  $\beta$  for each term.

```
beta_linear = np.random.uniform(-5, 5, size=n_features)

# Coefficients for interaction terms (n_features choose 2)
n_interactions = int(n_features * (n_features - 1) / 2)
beta_interaction = np.random.uniform(-1, 1, size=n_interactions)

# Linear combination of inputs
linear_terms = np.dot(X, beta_linear)

# Initialize interaction terms array
interaction_terms = np.zeros(n_samples)

# Index to keep track of interaction coefficients
idx = 0
```

```

# Loop over all unique pairs of features
for i in range(n_features):
    for j in range(i + 1, n_features):
        # Add interaction term
        interaction_terms += beta_interaction[idx] * X[:, i] * X[:, j]
        idx += 1

epsilon = np.random.normal(0, 1, n_samples)

# Total output
y = beta_linear[0] + linear_terms + interaction_terms + epsilon

# Feature names
columns = [f'x{i+1}' for i in range(n_features)]

# Combine features and target into a DataFrame
data = pd.DataFrame(X, columns=columns)
data['y'] = y

print(data.head(5))

```

	x1	x2	x3	x4	x5	x6
x7 \						
0	3.745401	9.507143	7.319939	5.986585	1.560186	1.559945
0.580836						
1	0.205845	9.699099	8.324426	2.123391	1.818250	1.834045
3.042422						
2	6.118529	1.394939	2.921446	3.663618	4.560700	7.851760
1.996738						
3	6.075449	1.705241	0.650516	9.488855	9.656320	8.083973
3.046138						
4	1.220382	4.951769	0.343885	9.093204	2.587800	6.625223
3.117111						

	x8	x9	x10	y
0	8.661761	6.011150	7.080726	178.654324
1	5.247564	4.319450	2.912291	85.918166
2	5.142344	5.924146	0.464504	76.613757
3	0.976721	6.842330	4.401525	-94.338920
4	5.200680	5.467103	1.848545	-38.892191

1. Data Preprocessing Handling Missing Values Our synthetic data has no missing values, but in real scenarios, you'd handle them here.

Feature Scaling Neural networks converge faster with scaled features.

```

from sklearn.preprocessing import StandardScaler

# Initialize scaler
scaler = StandardScaler()

```

```

# Fit and transform features
X_scaled = scaler.fit_transform(data.drop('y', axis=1))

# Update DataFrame
data_scaled = pd.DataFrame(X_scaled, columns=columns)
data_scaled['y'] = data['y']

```

Check Scaled Data

```
print(data_scaled.head())
```

	x1	x2	x3	x4	x5	x6
x7 \						
0	-0.363066	1.547478	0.821234	0.361551	-1.177298	-1.196034
1	-1.603520	1.612962	1.167403	-1.011189	-1.087152	-1.101643
2	0.468607	-1.219938	-0.694585	-0.463887	-0.129166	0.970644
3	0.453509	-1.114080	-1.477198	1.606041	1.650825	1.050610
4	-1.247970	-0.006553	-1.582870	1.465451	-0.818335	0.548268

	x8	x9	x10	y
0	1.302662	0.337660	0.761895	178.654324
1	0.121575	-0.252714	-0.690874	85.918166
2	0.085176	0.307297	-1.543968	76.613757
3	-1.355854	0.627727	-0.171851	-94.338920
4	0.105356	0.147797	-1.061607	-38.892191

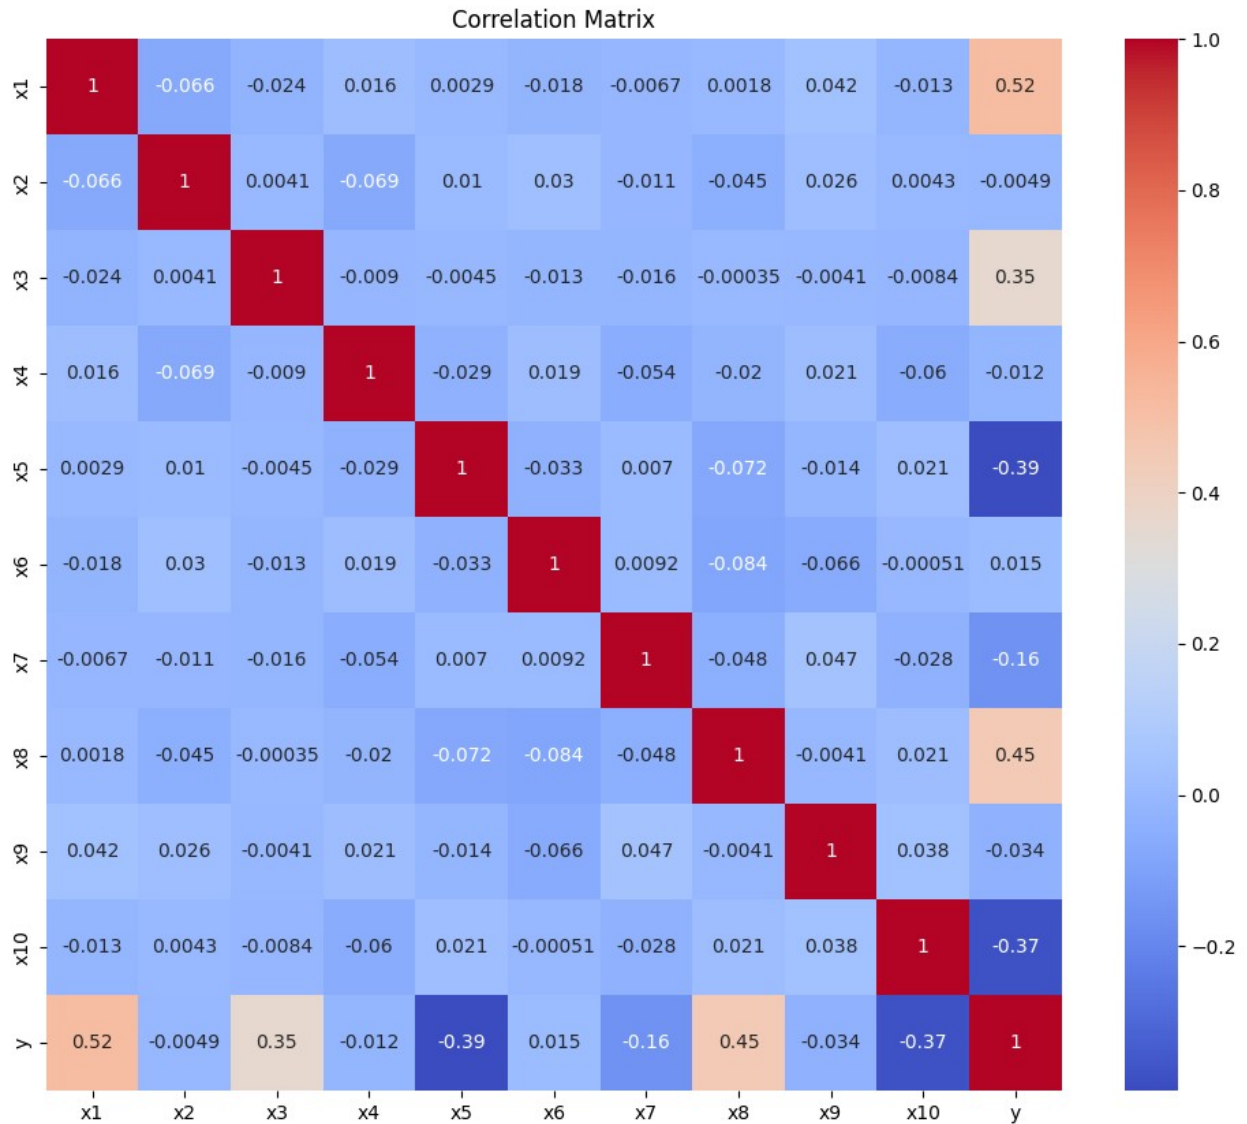
## 1. Feature Selection Correlation Analysis

```

# Compute correlation matrix
corr_matrix = data_scaled.corr()

# Plot heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm')
plt.title('Correlation Matrix')
plt.show()

```



Explanation:

Purpose: Identify highly correlated features. Action: We can decide to remove features with high multicollinearity. Select Features For simplicity, we'll proceed with all features, but in practice, you might select a subset.

1. Feature Engineering Creating Interaction Features Using DOE principles to include interaction terms.

```
from sklearn.preprocessing import PolynomialFeatures

# Create polynomial features (degree=2 includes interactions)
poly = PolynomialFeatures(degree=2, interaction_only=True,
include_bias=False)

# Fit and transform features
X_poly = poly.fit_transform(data_scaled.drop('y', axis=1))
```

```
# Get feature names
feature_names = poly.get_feature_names_out(columns)
```

Explanation PolynomialFeatures: degree=2: Includes all combinations of features up to degree 2.  
interaction\_only=True: Excludes squared terms, focusing on interactions.

Update DataFrame with Interaction Features

```
# Create DataFrame with interaction features
data_poly = pd.DataFrame(X_poly, columns=feature_names)
data_poly['y'] = data_scaled['y']
```

## 1. Neural Network Model Implementation Libraries

Split Data into Training and Testing Sets

```
from sklearn.model_selection import train_test_split

# Features and target variable
X = data_poly.drop('y', axis=1).values
y = data_poly['y'].values

# Split the data
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)

import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
```

Define the Neural Network Architecture

```
# Input dimension after adding interaction terms
input_dim = X_train.shape[1]

# Initialize the model
model = Sequential()

# Input layer
model.add(Dense(128, input_dim=input_dim, activation='relu'))

# Hidden layers
model.add(Dense(256, activation='relu'))
model.add(Dropout(0.3)) # Dropout for regularization
model.add(Dense(128, activation='relu'))

# Output layer
model.add(Dense(1, activation='linear'))
```

```
WARNING:tensorflow:From c:\Users\Yoga\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\backend.py:873: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.
```

```
model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 128)	7168
dense_1 (Dense)	(None, 256)	33024
dropout (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 128)	32896
dense_3 (Dense)	(None, 1)	129

```
=====  
Total params: 73217 (286.00 KB)
```

```
Trainable params: 73217 (286.00 KB)
```

```
Non-trainable params: 0 (0.00 Byte)
```

```
=====  
model.compile(  
    optimizer='adam',  
    loss='mean_squared_error',  
    metrics=['mean_absolute_error']  
)
```

```
WARNING:tensorflow:From c:\Users\Yoga\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\optimizers\__init__.py:309: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.
```

Explanation:

Optimizer: Adam optimizer adapts learning rates during training. Loss Function: Mean Squared Error for regression tasks. Metrics: Mean Absolute Error to evaluate model performance.

#### 1. Model Training Set Up Early Stopping

```
from tensorflow.keras.callbacks import EarlyStopping  
  
early_stop = EarlyStopping(  
    monitor='val_loss',
```

```
    patience=10,  
    restore_best_weights=True  
)
```

Train the Model

```
history = model.fit(  
    X_train, y_train,  
    validation_split=0.2,  
    epochs=100,  
    batch_size=32,  
    callbacks=[early_stop],  
    verbose=1  
)
```

Epoch 1/100

WARNING:tensorflow:From c:\Users\Yoga\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\utils\tf\_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From c:\Users\Yoga\AppData\Local\Programs\Python\Python39\lib\site-packages\keras\src\engine\base\_layer\_utils.py:384: The name tf.executing\_eagerly\_outside\_functions is deprecated. Please use tf.compat.v1.executing\_eagerly\_outside\_functions instead.

20/20 [=====] - 1s 9ms/step - loss: 8648.7988  
- mean\_absolute\_error: 73.9305 - val\_loss: 6188.2852 -  
val\_mean\_absolute\_error: 59.7385

Epoch 2/100

20/20 [=====] - 0s 3ms/step - loss: 8036.2563  
- mean\_absolute\_error: 71.0626 - val\_loss: 5332.3345 -  
val\_mean\_absolute\_error: 55.5473

Epoch 3/100

20/20 [=====] - 0s 3ms/step - loss: 4886.9434  
- mean\_absolute\_error: 53.6557 - val\_loss: 2312.6646 -  
val\_mean\_absolute\_error: 36.2888

Epoch 4/100

20/20 [=====] - 0s 3ms/step - loss: 1039.7753  
- mean\_absolute\_error: 24.2208 - val\_loss: 453.5901 -  
val\_mean\_absolute\_error: 16.8014

Epoch 5/100

20/20 [=====] - 0s 3ms/step - loss: 292.7521  
- mean\_absolute\_error: 13.1802 - val\_loss: 213.3192 -  
val\_mean\_absolute\_error: 11.3393

Epoch 6/100

20/20 [=====] - 0s 3ms/step - loss: 181.5477  
- mean\_absolute\_error: 10.2701 - val\_loss: 151.1905 -  
val\_mean\_absolute\_error: 9.8167



```
Epoch 7/100
20/20 [=====] - 0s 3ms/step - loss: 167.8574
- mean_absolute_error: 10.1185 - val_loss: 151.0729 -
val_mean_absolute_error: 9.5091
Epoch 8/100
20/20 [=====] - 0s 3ms/step - loss: 136.5610
- mean_absolute_error: 9.0322 - val_loss: 117.6982 -
val_mean_absolute_error: 8.3595
Epoch 9/100
20/20 [=====] - 0s 3ms/step - loss: 133.2912
- mean_absolute_error: 8.8254 - val_loss: 111.0519 -
val_mean_absolute_error: 7.9624
Epoch 10/100
20/20 [=====] - 0s 3ms/step - loss: 114.8808
- mean_absolute_error: 8.2795 - val_loss: 115.4543 -
val_mean_absolute_error: 8.3367
Epoch 11/100
20/20 [=====] - 0s 3ms/step - loss: 123.8664
- mean_absolute_error: 8.6454 - val_loss: 102.8765 -
val_mean_absolute_error: 7.8222
Epoch 12/100
20/20 [=====] - 0s 3ms/step - loss: 106.8634
- mean_absolute_error: 7.9053 - val_loss: 94.6192 -
val_mean_absolute_error: 7.4617
Epoch 13/100
20/20 [=====] - 0s 3ms/step - loss: 100.8406
- mean_absolute_error: 7.5977 - val_loss: 98.9632 -
val_mean_absolute_error: 7.5522
Epoch 14/100
20/20 [=====] - 0s 3ms/step - loss: 108.3899
- mean_absolute_error: 7.7403 - val_loss: 102.0158 -
val_mean_absolute_error: 7.5881
Epoch 15/100
20/20 [=====] - 0s 3ms/step - loss: 106.1804
- mean_absolute_error: 7.8147 - val_loss: 101.9140 -
val_mean_absolute_error: 7.7441
Epoch 16/100
20/20 [=====] - 0s 3ms/step - loss: 102.3587
- mean_absolute_error: 7.4590 - val_loss: 79.9470 -
val_mean_absolute_error: 6.7715
Epoch 17/100
20/20 [=====] - 0s 3ms/step - loss: 80.9630 -
mean_absolute_error: 6.8022 - val_loss: 81.3400 -
val_mean_absolute_error: 6.8737
Epoch 18/100
20/20 [=====] - 0s 3ms/step - loss: 84.8275 -
mean_absolute_error: 7.0133 - val_loss: 85.9735 -
val_mean_absolute_error: 7.0195
Epoch 19/100
```

```
20/20 [=====] - 0s 3ms/step - loss: 81.6997 -  
mean_absolute_error: 6.6208 - val_loss: 89.1710 -  
val_mean_absolute_error: 7.2989  
Epoch 20/100  
20/20 [=====] - 0s 3ms/step - loss: 80.7513 -  
mean_absolute_error: 6.7058 - val_loss: 84.8360 -  
val_mean_absolute_error: 7.0509  
Epoch 21/100  
20/20 [=====] - 0s 3ms/step - loss: 74.2774 -  
mean_absolute_error: 6.5526 - val_loss: 79.8263 -  
val_mean_absolute_error: 6.7771  
Epoch 22/100  
20/20 [=====] - 0s 3ms/step - loss: 72.4148 -  
mean_absolute_error: 6.4641 - val_loss: 79.8499 -  
val_mean_absolute_error: 6.8056  
Epoch 23/100  
20/20 [=====] - 0s 3ms/step - loss: 82.6289 -  
mean_absolute_error: 6.5561 - val_loss: 77.0856 -  
val_mean_absolute_error: 6.5755  
Epoch 24/100  
20/20 [=====] - 0s 3ms/step - loss: 67.6780 -  
mean_absolute_error: 6.1833 - val_loss: 80.7937 -  
val_mean_absolute_error: 6.7611  
Epoch 25/100  
20/20 [=====] - 0s 2ms/step - loss: 73.6766 -  
mean_absolute_error: 6.2056 - val_loss: 80.0893 -  
val_mean_absolute_error: 6.8974  
Epoch 26/100  
20/20 [=====] - 0s 3ms/step - loss: 71.0776 -  
mean_absolute_error: 6.3182 - val_loss: 66.9290 -  
val_mean_absolute_error: 6.4197  
Epoch 27/100  
20/20 [=====] - 0s 3ms/step - loss: 73.5757 -  
mean_absolute_error: 6.4198 - val_loss: 69.1794 -  
val_mean_absolute_error: 6.3932  
Epoch 28/100  
20/20 [=====] - 0s 3ms/step - loss: 79.4177 -  
mean_absolute_error: 6.4063 - val_loss: 77.1633 -  
val_mean_absolute_error: 6.6143  
Epoch 29/100  
20/20 [=====] - 0s 3ms/step - loss: 81.6471 -  
mean_absolute_error: 6.6154 - val_loss: 76.0644 -  
val_mean_absolute_error: 6.7377  
Epoch 30/100  
20/20 [=====] - 0s 2ms/step - loss: 66.4865 -  
mean_absolute_error: 6.1363 - val_loss: 78.9291 -  
val_mean_absolute_error: 6.9183  
Epoch 31/100  
20/20 [=====] - 0s 3ms/step - loss: 66.5892 -
```

```

mean_absolute_error: 6.1937 - val_loss: 66.3639 -
val_mean_absolute_error: 6.2720
Epoch 32/100
20/20 [=====] - 0s 2ms/step - loss: 60.3523 -
mean_absolute_error: 5.7764 - val_loss: 74.7099 -
val_mean_absolute_error: 6.6413
Epoch 33/100
20/20 [=====] - 0s 2ms/step - loss: 63.0094 -
mean_absolute_error: 6.0050 - val_loss: 75.5275 -
val_mean_absolute_error: 6.6400
Epoch 34/100
20/20 [=====] - 0s 2ms/step - loss: 61.7266 -
mean_absolute_error: 5.7045 - val_loss: 73.9673 -
val_mean_absolute_error: 6.5526
Epoch 35/100
20/20 [=====] - 0s 2ms/step - loss: 68.6619 -
mean_absolute_error: 6.1549 - val_loss: 77.3238 -
val_mean_absolute_error: 6.5474
Epoch 36/100
20/20 [=====] - 0s 2ms/step - loss: 62.0706 -
mean_absolute_error: 5.8968 - val_loss: 71.5218 -
val_mean_absolute_error: 6.5054
Epoch 37/100
20/20 [=====] - 0s 2ms/step - loss: 57.7333 -
mean_absolute_error: 5.6403 - val_loss: 70.3059 -
val_mean_absolute_error: 6.5438
Epoch 38/100
20/20 [=====] - 0s 3ms/step - loss: 50.6825 -
mean_absolute_error: 5.3124 - val_loss: 82.3516 -
val_mean_absolute_error: 7.0180
Epoch 39/100
20/20 [=====] - 0s 3ms/step - loss: 57.7500 -
mean_absolute_error: 5.6879 - val_loss: 81.5721 -
val_mean_absolute_error: 6.9372
Epoch 40/100
20/20 [=====] - 0s 3ms/step - loss: 57.5124 -
mean_absolute_error: 5.6813 - val_loss: 74.0030 -
val_mean_absolute_error: 6.8042
Epoch 41/100
20/20 [=====] - 0s 3ms/step - loss: 56.5801 -
mean_absolute_error: 5.6828 - val_loss: 68.1339 -
val_mean_absolute_error: 6.5257

```


Plot Training History

```

# Plot loss over epochs
plt.figure(figsize=(8, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')

```

```
plt.legend()
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Training History')
plt.show()
```



Explanation:

Early Stopping: Prevents overfitting by stopping training when validation loss doesn't improve.  
Batch Size: Set to 32 for efficient computation. Validation Split: 20% of training data used for validation

#### 1. Model Evaluation Evaluate on Test Data

```
test_loss, test_mae = model.evaluate(X_test, y_test, verbose=0)
print(f'Test MSE: {test_loss:.4f}')
print(f'Test MAE: {test_mae:.4f}')
```


Test MSE: 73.7218  
Test MAE: 6.5903

Predict on Test Data

```
y_pred = model.predict(X_test)
7/7 [=====] - 0s 1ms/step
from sklearn.metrics import r2_score, mean_squared_error
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse:.4f}')
print(f'R^2 Score: {r2:.4f}')
Mean Squared Error: 73.7218
R^2 Score: 0.9906
```

Plot Predicted vs Actual Values

```
plt.figure(figsize=(8, 6))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs Predicted Values')
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
         'k--', lw=2)
plt.show()
```



Explanation:

$R^2$  Score: Indicates how well the model explains the variability of the target variable. Mean Squared Error: Average squared difference between predicted and actual values.

1. Optimization and Conclusion Hyperparameter Tuning To improve the model, consider tuning hyperparameters:

Number of Layers and Neurons: Experiment with different architectures. Learning Rate: Adjust optimizer learning rate. Batch Size and Epochs: Modify based on convergence behavior. Activation Functions: Try different activation functions like tanh or leaky ReLU. Cross-Validation Implement k-fold cross-validation to ensure model robustness.

```
from sklearn.model_selection import KFold

# Define KFold
kfold = KFold(n_splits=5, shuffle=True, random_state=42)

# Placeholder for results
mse_scores = []
```

```

for train_index, val_index in kfold.split(X):
    # Split data
    X_train_fold, X_val_fold = X[train_index], X[val_index]
    y_train_fold, y_val_fold = y[train_index], y[val_index]

    # Build model (you might want to define a function for this)
    model_fold = Sequential()
    model_fold.add(Dense(64, input_dim=n_features, activation='relu'))
    model_fold.add(Dense(128, activation='relu'))
    model_fold.add(Dropout(0.2))
    model_fold.add(Dense(64, activation='relu'))
    model_fold.add(Dense(1, activation='linear'))
    model_fold.compile(optimizer='adam', loss='mean_squared_error')

    # Train model
    model_fold.fit(X_train_fold, y_train_fold, epochs=50,
batch_size=32, verbose=0)

    # Evaluate model
    y_val_pred = model_fold.predict(X_val_fold)
    mse_fold = mean_squared_error(y_val_fold, y_val_pred)
    mse_scores.append(mse_fold)


print(f'Cross-Validated MSE: {np.mean(mse_scores):.4f} ±
{np.std(mse_scores):.4f}')

7/7 [=====] - 0s 1ms/step
7/7 [=====] - 0s 1ms/step
7/7 [=====] - 0s 1ms/step
7/7 [=====] - 0s 1ms/step
7/7 [=====] - 0s 1000us/step
Cross-Validated MSE: 487.9427 ± 56.5823

# Calculate residuals
residuals = y_test - y_pred.flatten()

# Plot control chart
plt.figure(figsize=(12, 6))
plt.plot(residuals, marker='o', linestyle='-', color='b')
plt.axhline(y=np.mean(residuals), color='r', linestyle='--')
plt.axhline(y=np.mean(residuals) + 3*np.std(residuals), color='g',
linestyle='--')
plt.axhline(y=np.mean(residuals) - 3*np.std(residuals), color='g',
linestyle='--')
plt.title('Control Chart of Residuals')
plt.xlabel('Observation')
plt.ylabel('Residual')
plt.show()

```



Explanation:

Control Limits: Set at  $\pm 3$  standard deviations from the mean residual. Purpose: Detect any out-of-control signals indicating model issues. Conclusion Model Performance: Our neural network captures the complex relationships in the data. Optimization: Further tuning and feature engineering can enhance performance. SPC Application: Helps in monitoring and maintaining model reliability over time.


Next Steps Feature Engineering: Incorporate more interaction terms and nonlinear transformations. Advanced Techniques: Use models like XGBoost or Random Forests for comparison. Deployment: Implement the model in a production environment with continuous monitoring.

### 1. Applying Statistical Process Control Residual Analysis

```
# Calculate residuals
residuals = y_test - y_pred.flatten()

# Plot histogram of residuals
plt.figure(figsize=(8, 6))
sns.histplot(residuals, kde=True)
plt.title('Distribution of Residuals')
plt.xlabel('Residual')
plt.ylabel('Frequency')
plt.show()
```






Control Chart of Residuals

```
# Control chart parameters
mean_residual = np.mean(residuals)
std_residual = np.std(residuals)
upper_control_limit = mean_residual + 3 * std_residual
lower_control_limit = mean_residual - 3 * std_residual

# Plot control chart
plt.figure(figsize=(12, 6))
plt.plot(residuals, marker='o', linestyle='-', color='b',
label='Residuals')
plt.axhline(y=mean_residual, color='r', linestyle='--', label='Mean')
plt.axhline(y=upper_control_limit, color='g', linestyle='--',
label='Upper Control Limit')
plt.axhline(y=lower_control_limit, color='g', linestyle='--',
label='Lower Control Limit')
plt.title('Control Chart of Residuals')
plt.xlabel('Observation')
plt.ylabel('Residual')
```

```
plt.legend()
plt.show()
```



Explanation Purpose: Identify any out-of-control points indicating model issues. Control Limits: Set at  $\pm 3$  standard deviations from the mean residual.

### Process Capability Analysis

```
# Process capability indices
process_std = np.std(y_test)
spec_limits = [np.min(y_test), np.max(y_test)] # Assuming
specifications are the min and max of y_test
process_mean = np.mean(y_test)

# Calculate Cp and Cpk
Cp = (spec_limits[1] - spec_limits[0]) / (6 * process_std)
Cpk = min((spec_limits[1] - process_mean), (process_mean -
spec_limits[0])) / (3 * process_std)

print(f'Process Capability Cp: {Cp:.4f}')
print(f'Process Capability Cpk: {Cpk:.4f}')
```

```
Process Capability Cp: 1.0575
Process Capability Cpk: 0.9323
```

Explanation Cp and Cpk: Indices to assess the process capability in meeting specifications. Interpretation: Cp > 1: Process has the potential to meet specifications. Cpk > 1: Process is centered between the specification limits.

1. **Quality Management Considerations** Plan-Do-Check-Act (PDCA) Cycle Plan: Defined the problem, objectives, and planned data generation and modeling approach. Do: Executed data generation, preprocessing, and model training. Check: Evaluated model performance and applied SPC tools. Act: Identified areas for improvement, such as model tuning or data quality enhancements. Continuous Improvement Feedback Loop: Use evaluation results to refine the model. Employee Training: Ensure team members understand SPC and DOE principles applied in modeling. Documentation: Maintain detailed records of modeling procedures and findings. Compliance with ISO Standards ISO 9001 Principles: Customer Focus: Model aims to improve process quality, benefiting customers. Process Approach: Systematic modeling process aligns with process approach principle. Improvement: Continuous model evaluation and refinement.
2. **Conclusion** We successfully developed a neural network model to optimize a complex process by incorporating SPC and DOE principles. The model captures linear relationships and interactions between variables, providing valuable insights into the process dynamics.

**Key Takeaways** Data Generation: Simulated a realistic process with interactions and variability. Modeling Approach: Employed neural networks capable of handling high-dimensional data. Quality Tools: Applied SPC methods to monitor and control model performance. Quality Management: Aligned the project with quality management principles for continuous improvement.